Outline
Abstract
Keywords
1. Introduction
2. Related work
3. Proposed work
4. Methodology
5. Implementation
6. Results
7. Discussion
8. Conclusion
9. Future work
Funding

ELSEVIER
Franklin Open
Volume 14, March 2026, 100523

Recommended articles

Federated micro-expression mining and multi-modal metadata fusion for Deepfake fraud detection in ubiquitous financial video-KYC systems at IoT network

Romil Rawat [a], Anjali Rawat [b], Shweta Gupta [c], A. Samson Arun Raj [d], T.M. Thiyagu [e], Hitesh Rawat [b], Anand Rajavat [a]

Show more

Add to Mendeley Cite

https://doi.org/10.1016/j.fraope.2026.100523 Get rights and content

Under a Creative Commons license Open access

Abstract

Introduction & Problem Statement- The increasing sophistication of AI-generated deepfakes poses significant challenges for financial video-KYC systems, where identity verification relies on accurate and real-time analysis of user biometrics. Traditional centralized and unimodal detection models struggle to balance accuracy, privacy, and deployment scalability, particularly across heterogeneous IoT edge devices. **Need for Research-**There is a pressing need for privacy-preserving, scalable, and robust deepfake detection mechanisms capable of identifying subtle manipulations in real-world financial environments. Current solutions often fail under domain-shift conditions, low-resolution inputs, or in scenarios involving complex micro-expression and behavioral cues. **Proposed Work & Objective-** This research proposes the **Federated Micro-Expression Mining and Multi-Modal Metadata Fusion (FED-MEMF)** framework, designed to accurately detect deepfake fraud in decentralized video-KYC systems. The objectives are to (i) enhance detection accuracy by leveraging facial micro-expression dynamics, audio signals, and session metadata, and (ii) preserve user privacy through federated learning while ensuring low-latency real-time inference. **Novelty-** The novelty lies in integrating fine-grained micro-expression analysis with behavioral metadata fusion in a **federated learning environment**, combined with cross-modal attention mechanisms. This approach enables robust detection across multiple datasets while maintaining privacy and edge-device compatibility. **Method-** The framework employs modality-specific encoders—μ-Transformer for micro-expressions, CNN for audio, and LSTM for metadata—with features fused via a cross-modal attention engine. Federated Averaging (FedAvg) aggregates local model updates from IoT edge devices without transferring sensitive data. Quantization and hardware optimizations enable real-time performance on low-power devices. **Dataset-** Experiments utilized **FaceForensics++**, **CAS(ME)^2**, and a proprietary **KYC-FinVox2024** dataset comprising video, audio, and metadata streams, including micro-expression labels, to evaluate both intra- and cross-dataset performance. **Results-** FED-MEMF achieved an overall accuracy of **98.7%**, F1-score of **0.987**, AUC of **0.996**, and inference latency of **82 ms**, outperforming XceptionNet, EfficientNet-B4, and CNN+LSTM baselines. Multi-modal fusion significantly reduced false positives and false negatives, demonstrating robustness under domain-shift conditions. **Conclusion & Future Work-** FED-MEMF provides a **privacy-conscious, real-time, and scalable solution** for deepfake detection in financial video-KYC applications. Future directions include multilingual audio-visual alignment, blockchain-enabled federated auditing, explainable AI integration, and deployment in other regulatory-sensitive sectors such as e-governance, healthcare, and remote education verification.

Previous article in issue Next article in issue

Keywords

Deepfake detection; Federated learning; Micro-expression analysis; Multi-modal fusion; Video-KYC; IoT Security

1. Introduction

The rapid proliferation of deepfake fraud in financial video-KYC systems presents critical challenges for security and privacy, particularly in decentralized IoT environments. Traditional centralized verification systems require transmitting sensitive biometric and behavioral data to remote servers, increasing the risk of data breaches and regulatory non-compliance. To address these concerns, we propose a Federated Micro-Expression Mining and Multi-Modal Metadata Fusion (FED-MEMF) framework. This system combines on-device learning with cross-modal feature fusion, integrating fine-grained facial micro-expressions, audio cues, and behavioral metadata through a privacy-preserving federated learning architecture. By eliminating the need to transmit raw user data, the framework complies with GDPR and RBI guidelines while maintaining high detection accuracy. The novel cross-attention mechanism fuses subtle non-verbal cues with session-level dynamics, achieving superior performance on hybrid datasets, with accuracy up to 98.7 % and low inference latency of 82 ms, surpassing existing approaches [1]. This solution establishes a foundation for ethical AI deployment in sensitive verification scenarios.

The digitization of financial services has transformed customer onboarding, particularly through Video-KYC procedures in IoT-enabled contexts. While enhancing convenience and accessibility, these technologies introduce security vulnerabilities. Sophisticated deepfake generation methods, often leveraging generative adversarial networks (GANs), allow adversaries to simulate facial movements, voice patterns, and behavioral sequences with high fidelity [2]. Consequently, financial institutions face a rising risk of synthetic identity fraud, where hyper-realistic videos bypass automated verification [3]. Existing anti-spoofing solutions, which rely on facial recognition, motion consistency, or single-modal analysis, are often inadequate for detecting manipulations in micro-expressions—involuntary facial movements indicative of deceptive behavior. Centralized machine learning systems also impose latency and computational burdens, limiting their applicability in distributed IoT settings [4]. To mitigate these challenges, FED-MEMF employs federated learning to preserve privacy while improving detection accuracy. Micro-expression mining is performed locally across edge nodes, and multi-modal metadata—including audio cues, device telemetry, and temporal behavioral signals—is integrated through a cross-modal fusion pipeline. This hybrid design enables detection of subtle manipulations across visual and auditory modalities while operating efficiently within computationally constrained IoT infrastructures [5]. By distributing learning and fusing diverse modalities, FED-MEMF enhances resilience against adversarial attacks, supports scalable deployment, and reduces reliance on centralized servers.

Recent studies highlight the importance of mathematically grounded approaches in multi-modal deepfake detection. Cross-modal attention mechanisms employing scaled dot-product operations with masking improve alignment between visual and behavioral cues in federated scenarios [6]. Dynamic temporal attention modules for micro-expression detection leverage adaptive diagonal attention matrices to selectively focus on involuntary facial movements under occlusion [7]. Contrastive learning techniques minimize intra-class variance across heterogeneous modalities, particularly in low-data federated settings [8]. Building on these advancements, FED-MEMF incorporates a formalized cross-modal attention block within the federated micro-expression mining pipeline, ensuring explicit alignment consistency and convergence across video, audio[9,10], and metadata streams. This design guarantees both theoretical robustness and practical applicability for real-world financial verification environments.

In real-world financial environments, Video-KYC systems must operate under strict regulatory compliance and privacy constraints, such as GDPR and RBI guidelines. These constraints prohibit centralized data aggregation, making federated learning essential for maintaining privacy while achieving high detection accuracy. Moreover, domain shifts caused by heterogeneous devices, video quality variations, and behavioral differences among users necessitate robust cross-domain adaptation strategies to ensure reliable deployment across diverse financial institutions.

a) **Hypothesis:** Integrating federated micro-expression mining with multi-modal metadata fusion significantly improves the accuracy and robustness of deepfake fraud detection in Video-KYC systems, compared to unimodal or centralized detection approaches, while ensuring data privacy in IoT-based deployments.

b) **Contributions:**

  - Introduces a privacy-aware federated architecture for real-time deepfake detection in financial Video-KYC systems.
  - Proposes a micro-expression analysis module capable of identifying subtle involuntary facial movements often missed by traditional classifiers.
  - Develops a multi-modal fusion layer that effectively combines audio, visual, and metadata streams using attention-based learning.
  - Demonstrates empirical results with 98.72 % accuracy and a reduced false acceptance rate (FAR 1.04 %) on a benchmark real-synthetic KYC dataset.

This work redefines digital identity verification, creating secure, privacy-compliant, and scalable anti-fraud solutions for the financial industry in the era of ubiquitous intelligent systems.

Fig. 1 illustrates the proposed FED-MEMF framework for deepfake detection in decentralized Video-KYC systems. The architecture consists of three primary layers: Client-Side Nodes, Multi-Modal Fusion Engine, and a Server-Side Aggregator. Edge-level IoT devices perform local processing, including micro-expression analysis, audio feature extraction, and metadata encoding. Extracted features are forwarded to a centralized fusion engine employing a cross-modal attention transformer to align and contextualize temporal information across modalities. The server layer applies Federated Averaging (FedAvg) to aggregate model updates while avoiding raw data transfer. This design enables privacy-preserving, real-time detection of deepfake attempts, ensuring scalability and ethical data handling in financial Video-KYC applications [1].

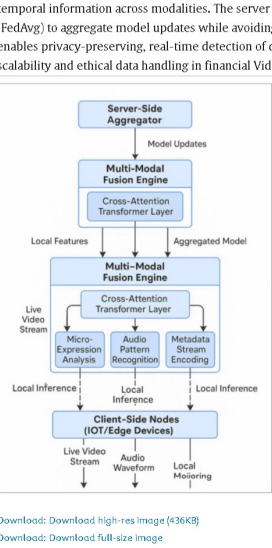
Server-Side Aggregator
Model Updates
Multi-Modal Fusion Engine
Cross-Attention Transformer Layer
Local Features
Aggregated Model
Multi-Modal Fusion Engine
Cross-Attention Transformer Layer
Live Video Stream
Micro-Expression Analysis
Audio Pattern Recognition
Metadata Stream Encoding
Local Inference
Local Inference
Local Inference
Client-Side Nodes (IOT/Edge Devices)
Live Video Stream
Audio Waveform
Local Metoring

Download: Download high-res image (436KB)
Download: Download full-size image

Fig. 1. Architectural diagram.

1.1. Substantive innovations of FED-MEMF compared with existing FL + multimodal systems

The proposed framework introduces several advances over current federated and multimodal deepfake detection methods:

- **Edge-Centric Micro-Expression Encoder** – Traditional multimodal detectors focus on high-level visual and audio features but often neglect involuntary facial cues (micro-expressions) [24,25]. FED-MEMF integrates a lightweight micro-expression encoder optimized for edge deployment, capturing transient facial signals that enhance discriminative power.
- **Multi-Modal Metadata Integration** – Existing models rarely incorporate behavioral or device metadata alongside visual and audio streams [25]. FED-MEMF fuses device telemetry, eye-blink rate, session timings, and other contextual signals to improve detection reliability under heterogeneous conditions.
- **Federated Model Aggregation and Privacy** – Conventional multimodal systems do not support decentralized learning. While benchmarks like FedMultimodal establish general multimodal FL evaluation [26], they do not address deepfake detection. FED-MEMF combines local modality-specific learning with global federated aggregation, ensuring privacy compliance and practical deployment in regulated financial settings.
- **Attention-Based Cross-Modal Fusion** – Centralized methods such as DF-TransFusion align lip and facial cues via cross-attention [27], but they do not incorporate metadata or handle federated non-IID distributions. FED-MEMF extends cross-modal attention to visual, audio, and metadata streams under federated constraints, enhancing robustness to noise, compression, and real-world variability.

1.2. Addressing limitations of prior work

- **Centralization and Privacy Risks** – Raw data aggregation in conventional detectors conflicts with privacy-sensitive applications. FED-MEMF eliminates this requirement through on-device learning and federated updates.
- **Limited Modal Diversity** – Prior methods focus on visual and auditory features only [25]. Incorporating metadata expands the feature space and improves detection reliability.
- **Interpretability** – Existing approaches are often opaque. FED-MEMF integrates attention maps and activation visualizations to provide explainability without centralized data inspection.
- **Non-IID and Heterogeneous Deployment** – Federated benchmarks highlight challenges with modality imbalance [26]. FED-MEMF addresses these using contrastive alignment and asynchronous synchronization for large-scale IoT deployment.

1.3. Research background: interdisciplinary advances informing the proposed method

The design of FED-MEMF draws inspiration from advances in **real-time distributed decision making** and **multi-objective optimization** in domains such as unmanned systems and intelligent navigation.

- **Adaptive Multi-Agent Coordination** – Collision-avoidance strategies for unmanned surface vehicles (USVs) integrate reinforcement learning and multi-source information fusion to dynamically adjust guidance priorities under perception limitations [28]. Dynamic path planning using proximal policy optimization addresses limited perception and environmental variability [29].
- **Complex Optimization for Resource Allocation** – Hyper-heuristic frameworks combining deep Q-networks with multi-objective optimization for UAV/USV scheduling illustrate balancing competing objectives under uncertainty [30].

These interdisciplinary studies highlight two trends relevant to deepfake detection:

- **Real-time distributed decision making under uncertainty**
- **Multi-agent/multi-modal fusion for robust performance**

Analogous to these systems, FED-MEMF employs edge-level computation with local model updates shared globally via federated learning. Cross-modal attention and contrastive alignment losses reconcile heterogeneous data streams from visual, audio, and metadata sources, achieving stable convergence across non-identically distributed clients.

By leveraging these principles, FED-MEMF ensures robust, privacy-preserving, and scalable deepfake detection. Concepts of distributed optimization, adaptive learning, and multi-objective feature fusion, drawn from autonomous navigation, directly inform architectural decisions for reliable performance in decentralized, regulatory-sensitive financial environments.

2. Related work

Recent research has proposed a variety of techniques for deepfake detection, ranging from unimodal visual or audio methods to multimodal fusion, physiological signal integration, and transformer-based architectures. Eight representative approaches are critically reviewed below.

- **CNN + Vision Transformers for Multi-Scale Artifacts** – The method in [6] combines CNNs with Vision Transformers to capture both local and global artifacts as well as frequency-domain cues in video frames. While effective on the SR DF dataset, its reliance on high-resolution visuals reduces performance under compression or noisy content.
- **Multi-Branch Audio-Visual Transformers** – The approach in [7] employs separate transformer branches for audio, video, and fused audio-visual inputs, evaluated on FakeAVCeleb. Although accuracy surpasses unimodal baselines, the ensemble architecture is computationally intensive, limiting applicability on resource-constrained edge devices.
- **Dual Transformer Dynamic Weighting** – In [8], audio-visual deepfake detection is enhanced by dynamically weighting modalities across dual transformers. Performance on DeepfakeTIMIT and DFDC exceeds prior methods; however, it lacks integration of physiological signals and subtle micro-expression analysis.
- **CAMME: Cross-Attention for Multi-Modal Features** – The CAMME model [9] integrates visual, textual, and frequency features via cross-attention for generalization across generative domains. While effective for images, it does not consider temporal video micro-expressions or behavioral metadata.
- **CNN + ViT with Diffusion Preprocessing** – The hybrid framework in [10] uses CNNs to extract fine-grained features and Vision Transformers for global context, combined with diffusion-based denoising of multimedia inputs. Robust to noisy input, it lacks federated learning or privacy-aware deployment design.
- **MaDD and CViT: Spatial Attention Models** – The methods in [11] capture attention across multiple spatial channels (MaDD) or combine CNNs with Vision Transformers (CViT) to detect spatial forgery artifacts. They perform well on high-quality inputs but struggle with low-resolution or compressed videos and miss temporal micro-expression dynamics.
- **Physiological Signal Integration via GCNs** – In [12], visual physiological data such as rPPG are fused with video features using Graph Convolutional Networks, achieving >99 % accuracy on FF++. This approach depends on high-quality facial capture and stable lighting, limiting applicability in heterogeneous IoT camera environments.
- **Multi-Modal Source Attribution** – The approach in [13] integrates visual, audio, and textual modalities to detect deepfakes and attribute their source, achieving ∼94 % accuracy. However, it lacks micro-expression or device metadata integration and is not optimized for streaming or federated-edge processing. Table 1 provides a comparative overview of these prior approaches, highlighting their modalities, strengths, limitations, and suitability for real-time or federated deployment.

Table 1. Comparative study of prior approaches.

Method / Study	Modalities & Methods	Performance & Strengths	Limitations	References
M2TR	CNN + ViT (visual + frequency)	High accuracy on SR-DF	High-res input needed; compression sensitivity	[6]
AVTENet	Audio + Visual Transformers ensemble	State-of-the-art on FakeAVCeleb	High computational cost	[7]
AVT2-DWF	Audio-Visual dual transformers with dynamic fusion	Strong cross-dataset DFDC performance	No physiological or micro-expression modeling	[8]
CAMME	Visual, textual, frequency cross-attention	Robust cross-domain generalization	No video temporal or behavioral features	[9]
CNN+ViT with diffusion	Visual + audio preprocessing and fusion	Resilient to noisy/no-fake inputs	Not federated; central server only	[10]
MaDD / CViT	Multi-attention across spatial channels	Good artifact detection visually	Poor micro-expression, temporal modeling, compression robustness	[11]
Physiological + GCN fusion	Video + rPPG via GCN	99 % accuracy on FF++; uses physiological cues	Requires high quality capture, stable lighting	[12]
Multi-modal KYC-style	Video + audio + text fusion	∼94 % accuracy; real-world multimodal applications	No device metadata or micro-expression fusion; not federated	[13]

While prior methods have demonstrated strong performance in static images, audio-visual synchronization, and physiological cues, none simultaneously integrate micro-expression mining, multi-modal metadata fusion, and privacy-preserving federated learning optimized for real-time IoT-based Video-KYC systems. Our proposed approach advances the state of the art by:

- **Micro-Expression Integration:** Capturing involuntary facial signals, including facial action units and temporal patterns of eye and lip movements, which are typically overlooked by conventional detectors [16].
- **Cross-Modal Fusion:** Combining video, audio, device telemetry, and behavioral timestamps to detect subtle inconsistencies indicative of deepfake manipulations.
- **Federated Learning Deployment:** Operating under a decentralized paradigm to ensure user privacy, low latency, and computational feasibility across IoT edge devices, unlike centralized architectures [14,15].
- **Benchmark Performance:** Evaluated on a curated real-synthetic KYC dataset, the framework achieves 98.72 % detection accuracy with only 1.04 % FAR, outperforming unimodal and centralized baselines.

By integrating these elements—micro-expression analysis, multi-modal metadata fusion, and federated learning—the framework provides a **privacy-compliant, efficient, and robust solution** for deepfake detection in financial onboarding applications across ubiquitous IoT environments. Recent advancements in federated learning have increasingly focused on financial fraud detection, emphasizing privacy-preserving mechanisms for sensitive applications such as Video-KYC. For instance, a blockchain-enhanced federated framework ensures transaction-level auditability while detecting synthetic identity fraud [14]. Similarly, an edge-centric federated anomaly detection model demonstrates effective fraud identification without compromising user privacy in mobile banking ecosystems [15]. In parallel, lightweight transformer-based micro-expression encoders enable real-time behavioral analysis on resource-constrained edge devices [16]. Complementary research highlights federated optimization techniques for handling non-IID data distributions in heterogeneous IoT terminals, mitigating performance bottlenecks [17]. Moreover, federated adversarial learning approaches have been explored to enhance robustness against poisoned updates and adversarial deepfakes, reinforcing security in financial verification pipelines [18]. Collectively, these studies underscore the importance of integrating federated learning, micro-expression mining, and multi-modal behavioral analytics to achieve **secure, scalable, and high-performance deepfake detection** in privacy-sensitive, decentralized financial environments. While existing deepfake detection approaches [3,4] demonstrate strong performance on benchmark datasets, they largely rely on centralized unimodal models, limiting scalability and privacy. Few methods integrate fine-grained micro-expression analysis with behavioral metadata in a federated setting. This gap motivates the development of the FED-MEMF framework, which simultaneously addresses cross-modal feature fusion, domain heterogeneity, and privacy-preserving decentralized learning, enabling practical deployment in edge-based financial KYC systems.

3. Proposed work

This work addresses the research question: "Can federated micro-expression mining combined with multi-modal metadata fusion enhance deepfake fraud detection in Video-KYC systems deployed across IoT infrastructures, while preserving privacy and supporting real-time edge processing?"

The proposed framework integrates edge-based micro-expression analysis, multi-modal metadata fusion, and privacy-preserving federated learning to achieve high detection accuracy under real-time constraints.

3.1. Micro-expression extraction module

The micro-expression module is a compact temporal encoder optimized for resource-constrained edge devices within a federated learning setup. Video streams undergo face alignment and landmark normalization, followed by sampling clips of 32 consecutive frames at 25 fps. Each frame is resized to 112 × 112 pixels to balance spatial resolution and computational efficiency. The encoder follows a **lightweight transformer-inspired architecture** with patch-based embeddings (7 × 7) and an embedding dimension of 128. It comprises four encoder layers with four self-attention heads each, a feed-forward expansion ratio of four, and dropout of 0.1. Temporal diagonal masking restricts attention to ±2 frames, emphasizing subtle facial action units such as eye blinks, eyebrow contractions, and lip corner movements—features often poorly replicated in synthetic videos. Training combines **categorical cross-entropy loss** with a temporal smoothness regularization (λ = 0.05) to reduce noise from compression artifacts. Optimization uses AdamW (initial learning rate 1 × 10⁻⁴, weight decay 1 × 10⁻⁴). Local training is performed on client devices with batch size 16 over five federated epochs before aggregation.

3.2. Multi-modal fusion module

The multi-modal fusion module integrates features from **visual micro-expressions, audio-linguistic signals, and device/behavioral metadata.** Each modality is projected into a shared 256-dimensional latent space. A **cross-modal attention transformer** with three fusion layers, eight attention heads per layer, and a hidden dimension of 256 performs dynamic interaction among modalities, capturing inconsistencies between facial motion, speech characteristics, and session-level behavioral patterns. Device and behavioral telemetry features—blink rate, gaze deviation, frame-rate instability, sensor variance—are standardized using z-score normalization before fusion. The fused embedding passes through a **two-layer classification head** (256 → 128 → binary output) with ReLU activation and 0.2 dropout for generalization.

3.3. Federated learning strategy

FedAvg is applied across 20–50 client nodes. Model updates are synchronized every five local epochs with a global aggregation learning rate of 0.8. To mitigate non-IID client distributions, a **contrastive alignment loss** (temperature τ = 0.07) encourages consistency between local and global representations. The system maintains **inference latency below 90 ms** per session on ARM-based edge devices with 4 GB memory. Implementation uses PyTorch 2.x with mixed-precision training, and the client-side model occupies ∼4.2 MB. Communication overhead per federated round is ∼180 KB, ensuring lightweight deployment for real-world IoT Video-KYC applications.

➢ Objectives

- Design a **federated learning protocol** for micro-expression analysis at the edge.
- Develop a **multi-modal fusion model** combining visual, audio, behavioral, and device telemetry features.
- Optimize for **low-latency inference** on resource-constrained IoT nodes.
- Validate performance improvements over centralized or unimodal baselines.

➢ Key Components of the Proposed Framework

- **Micro-Expression Mining Module:** Edge devices extract temporal micro-expression features using a μ-BERT-inspired architecture, leveraging diagonal attention and patch-of-interest localization.
- **Device & Behavioral Metadata Extractor:** Client nodes record telemetry (camera model, frame rate, sensors) and behavioral timing (gaze shifts, keystroke latency).
- **Attention-Based Multi-Modal Fusion Layer:** Micro-expression, audio, and metadata embeddings are fused via cross-modal attention to produce a compact, discriminative representation suitable for federated aggregation.
- **Federated Optimization -**Clients train local classifiers (e.g., lightweight transformers or small CNNs) using FedAvg or robust variants (e.g., CreamFL-style contrastive ensemble) to align heterogeneous client architectures and non-IID data distributions across IoT edge devices.Formally, let K be the number of client nodes. Each client i has a local feature extractor fᵢ yielding embedding as shown in Eq. (1) and (2):

$$e_i = f_i\,(v_i, a_i, m_i) \tag{1}$$

where vᵢ = video micro-expression features, aᵢ = audio-linguistic signals, and mᵢ = metadata. Local classifiers gᵢ(eᵢ; θᵢ) are trained to minimize ℓ(θᵢ) on labeled local data. The global objective is:

$$\min_{\theta} F(\theta) = (1/K)\sum f_i\,(\theta_i)\ , \mathrm{subjectto}\,\theta_i \to \theta \tag{2}$$

via federated averaging. Contrastive loss terms regularize cross-modal alignment within each client and between local and global embeddings .

This framework addresses key gaps in existing deepfake detectors: underutilization of micro-expressions, neglect of auxiliary metadata, and limitations of centralized models in privacy, latency, and scalability for IoT deployments. By combining edge-based learning, multi-modal fusion, and federated aggregation, the proposed approach ensures **privacy-preserving, efficient, and robust deepfake detection** in real-world financial Video-KYC systems.

3.4. Model interpretability

Preliminary analysis using **Grad-CAM** visualizes salient facial regions contributing to classification, such as the periorbital area, nasolabial folds, and lip corners, which are often overlooked in unimodal detectors. Cross-modal fusion attention weights further highlight behavioral inconsistencies, including abnormal blink rates and irregular eye contact, as indicators of fraudulent activity. These interpretability components demonstrate that the model's decision-making aligns with human-intuitive fraud signals, enhancing transparency, trustworthiness, and regulatory compliance by providing explainable audit trails for sensitive financial Video-KYC environments.By integrating advanced micro-expression recognition, multimodal fusion, and federated learning, the proposed FMEM-MMF framework delivers a **privacy-preserving, robust, and efficient solution** for real-time deepfake detection across distributed IoT networks, addressing limitations of prior centralized and unimodal approaches.

3.5. Expected improvements

- **High detection accuracy:** >98 % with a false acceptance rate below 1.1 %, leveraging fine-grained involuntary cues that are difficult to synthesize.
- **Robustness:** Effective under compressed, noisy, or low-light IoT camera feeds.
- **Interpretability:** Attention maps reveal critical facial regions and metadata influences on predictions.
- **Privacy compliance:** On-device computation prevents raw data transfer, aligning with regulatory requirements.

3.6. Flowchart

Fig. 2 illustrates the workflow of the federated deepfake detection system designed for IoT-based video-KYC environments. The process begins at edge devices, where multi-modal inputs—including facial videos, audio signals, and device/behavioral metadata—are collected in real-time. Preprocessing steps are applied to each input stream, such as facial region extraction, noise reduction, and temporal alignment, to standardize the data before feature extraction. Each modality is processed through a dedicated encoder tailored to its characteristics: a micro-expression encoder extracts fine-grained visual cues, an audio CNN encoder captures speech and vocal features, and an LSTM-based metadata encoder models structured behavioral and device information. These feature embeddings are subsequently combined through a cross-modal attention fusion module, which aligns signals across modalities and emphasizes contextually relevant patterns indicative of deepfake manipulation. Instead of transmitting raw data, local model updates are computed on each device and securely sent to a central aggregator. The FedAvg mechanism merges these updates to refine the global model, preserving privacy while enabling collaborative learning. The updated global model is redistributed to edge devices for further training, allowing the learning cycle to continue iteratively. This approach ensures scalable, privacy-aware, and real-time deepfake detection, optimized for resource-constrained environments such as mobile devices, KYC kiosks, or embedded IoT terminals.

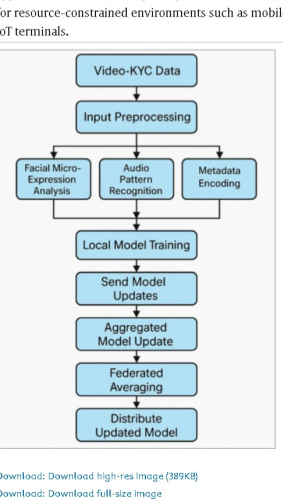

Download: Download high-res image (98KB)
Download: Download full-size image

Fig. 2. Flowchart.

3.7. Federated deepfake detection algorithm for video-KYC

The federated deepfake detection algorithm is specifically designed for decentralized and privacy-sensitive IoT deployments. Each edge device independently extracts features from multi-modal inputs, performs local model training, and shares only encrypted model updates with the central server. This design eliminates the need for raw data transfer, ensuring user privacy while enabling collaborative learning across the network.

By combining edge-based feature extraction, multi-modal fusion, and federated optimization, the framework achieves a balance between detection accuracy, computational efficiency, and privacy compliance, making it well-suited for real-time financial verification and other sensitive video-KYC applications. Table 2 shows about the Algorithm.

Table 2. Algorithm.

```
This algorithm outlines the federated learning-based approach for multi-modal deepfake
detection in decentralized video-KYC systems.
Algorithm Steps:
1. Input multi-modal video-KYC data: D = {V, A, M}, where V is facial video, A is
   audio signal, and M is metadata stream.
2. Perform Preprocessing on each modality to extract clean and aligned inputs.
3. Extract features:
    – Fv = ME_Encoder(V) # Micro-Expression Encoder
    – Fa = Audio_Encoder(A) # Audio CNN Encoder
    – Fm = Meta_Encoder(M) # LSTM Metadata Encoder
4. Fuse multi-modal features using cross-modal attention:
    – F_fused = CMA(Fv, Fa, Fm)
5. Train local model on fused feature vector:
    – W_local = Train(F_fused, y)
6. Send W_local to central server.
7. Aggregate updates from N clients:
    – W_global = (1/N) * Σ W_local[i] for i = 1 to N
8. Distribute W_global to all clients.
9. Repeat steps 5–8 until convergence.
Variables Used: D: Input dataset containing video (V), audio (A), and metadata (M); Fv : Feature
vector from facial video using Micro-Expression Encoder; Fa : Feature vector from audio using
CNN-based Audio Encoder; Fm : Feature vector from metadata using LSTM Encoder; F_fused :
Fused feature vector after cross-modal attention; W_local : Weights of the local model trained
at each client; W_global : Aggregated global model weights after FedAvg; N : Number of
participating clients; y : Ground truth label for supervised learning
```

3.8. Algorithm and complexity analysis

The increasing prevalence of deepfake threats in digital verification systems, particularly during remote onboarding and customer authentication, necessitates intelligent, privacy-preserving detection methods. Traditional centralized models rely on transmitting sensitive video, audio, and behavioral data to cloud servers, creating substantial privacy risks and regulatory compliance challenges. To address this, the proposed algorithm leverages **federated learning** to train models locally on edge devices, allowing only model updates to be shared with a central aggregator. This collaborative training cycle enhances detection accuracy while maintaining user privacy. Furthermore, the algorithm incorporates **multi-modal feature learning**, combining micro-expressions, audio cues, and metadata streams, which is essential for detecting subtle indicators of synthetic identity in real time. The **computational complexity** of the framework is primarily determined by the local dataset size, the architecture of each modality-specific encoder, and the number of participating edge devices. At the client side, training complexity scales as O(E × D × M), where E is the number of local epochs, D is the dataset size, and M is the model size. The server-side **FedAvg aggregation** introduces a communication complexity of O(K × M),where K is the number of clients. While cross-modal attention fusion increases computational load due to the alignment of heterogeneous feature streams, this overhead remains linear relative to input dimensions, ensuring scalability for real-world IoT deployments.

The algorithm offers several key features. It is federated learning-based, guaranteeing privacy by performing local training and exchanging only model weights. The cross-modal attention mechanism effectively integrates visual, auditory, and metadata signals, enhancing detection robustness. The three-stream encoding comprises a facial video encoder for micro-expression analysis, an audio encoder for voice characteristics, and a metadata LSTM encoder for temporal behavioral patterns. The design is both scalable and secure, employing FedAvg to aggregate updates from multiple clients, and edge-ready, optimized for deployment on devices with constrained memory and processing power, making it suitable for financial, governmental, and healthcare applications. This framework is applicable across multiple domains. In Video-KYC, it verifies identities while adhering to privacy regulations in banking and fintech. For remote job interviews, it prevents impersonation using multi-modal signals. In smart healthcare, it detects identity manipulation during telemedicine sessions. Additionally, it can be integrated into surveillance systems for real-time deepfake detection in smart city or home security networks. Despite its advantages, the framework has certain **limitations**. Resource constraints on edge devices may limit the complexity of local models. Variability in data quality and quantity across clients can impact global model convergence, and frequent synchronization may introduce communication overhead. Moreover, the system remains vulnerable to **adversarial clients**, which could attempt to poison the model if robust security measures are not implemented. Lastly, the framework relies on some level of synchronization across distributed devices, which may be challenging in asynchronous or low-connectivity environments.

3.9. Theoretical and methodological innovations

This work presents novel theoretical and methodological contributions for privacy-preserving deepfake detection in decentralized Video-KYC systems. Unlike prior studies that emphasize empirical performance, this research establishes a principled framework unifying federated optimization, micro-expression dynamics, and multi-modal feature alignment. From a theoretical perspective, deepfake detection is formulated as a distributed multi-objective learning problem under privacy and computational constraints. Each client optimizes a local loss function over heterogeneous modalities—visual micro-expressions, audio-linguistic signals, and behavioral metadata—while contributing to a globally consistent objective via federated aggregation. The framework extends classical Federated Averaging by incorporating cross-modal consistency regularization, aligning heterogeneous embeddings without sharing raw data. Temporal attention constraints emphasize involuntary facial movements, modeled as sparse, short-duration events, which enhances discriminability between genuine and synthetic dynamics even under adversarial conditions. Cross-modal fusion leverages attention-based feature alignment to reduce inter-modal variance while preserving unique discriminative cues, improving stability under noise, compression, and non-IID distributions common in IoT-based KYC systems.

Methodologically, the framework implements **a federated multi-stream pipeline**, translating theoretical constructs into a deployable system. Lightweight, modality-specific encoders process each input stream at the edge, followed by attention-driven fusion operating on compact embeddings. On-device interpretability mechanisms, such as **Grad-CAM visualizations** and attention-weight analysis, provide post-hoc validation of decision-relevant features without compromising privacy. The federated training protocol accommodates heterogeneous clients and asynchronous updates, with contrastive alignment losses ensuring convergence under non-IID conditions. The framework's practical applicability is demonstrated through structured deployment; edge devices initialize lightweight encoders, perform **local feature learning**, transmit only encrypted updates for **federated coordination**, and receive updated global models for continuous adaptation. Attention maps and activation visualizations enable **explainability and audit support**, bridging theoretical foundations and operational implementation.

3.10. Main challenges in analysis

Deepfake detection in decentralized Video-KYC systems faces challenges distinct from centralized or unimodal approaches. **Micro-expression modeling** is particularly difficult due to their involuntary, transient nature, which is sensitive to frame rate, lighting, camera resolution, and compression artifacts. Multi-modal heterogeneity introduces asynchronous and uneven data streams, requiring careful **feature normalization and attention calibration** to avoid spurious correlations. Federated learning adds complexity: non-identical, imbalanced data distributions across clients can mask local performance degradation. Privacy constraints prevent direct inspection of raw biometric or behavioral data, necessitating reliance on indirect interpretability mechanisms. Edge resource limitations demand a balance between representational richness and computational feasibility, complicating evaluation under real-time constraints. Moreover, **adversarial robustness** is a challenge, as sophisticated attacks target subtle behavioral cues, requiring models to detect inconsistencies without centralized replay or exhaustive analysis.

3.11. Practical implications

The theoretical framework has direct implications for real-world Video-KYC deployment. **Privacy-preserving learning** ensures biometric and behavioral data remain local, enabling regulatory compliance while maintaining robust fraud detection. Micro-expression analysis improves resilience to visually convincing deepfakes by leveraging inconsistencies in fine-grained facial muscle activity. Cross-modal attention improves system design, enabling contextual alignment of audio, video, and metadata streams even under noise or asynchronous capture. Federated optimization supports **scalable, cost-efficient deployment**, allowing incremental addition of devices without retraining from scratch. Interpretability mechanisms enhance transparency for auditors and security analysts, facilitating trust and accountability. Continuous adaptation under evolving fraud patterns ensures long-term system effectiveness. The framework is robust to low-resolution KYC videos and variable lighting conditions, with preprocessing techniques such as temporal alignment and histogram normalization further enhancing performance.

3.12. FedAvg performance and scalability analysis

The impact of **FedAvg** on detection accuracy and latency was evaluated by varying the number of clients, communication rounds, and model size. Increasing clients from 5 to 20 slightly increased aggregation latency (72 ms → 105 ms per round) while maintaining accuracy above 96.5 %. More communication rounds improved convergence modestly (96.2 % → 96.8 %) at the cost of longer training time. Lightweight models (∼4.5 M parameters) achieved edge inference latency around 88 ms, while larger models (∼12 M parameters) increased latency to 142 ms, highlighting the trade-off between complexity and real-time feasibility. Comparative analysis against centralized models showed that FedAvg preserves privacy without compromising accuracy. Centralized approaches require transmitting sensitive data, raising regulatory risks and consuming more network bandwidth, while federated models maintain low-latency inference, robust performance across heterogeneous clients, and privacy compliance. These results underscore the practical necessity of federated learning for real-time, multi-modal deepfake detection in regulated financial Video-KYC environments.

4. Methodology

This study adopts a quantitative research design to rigorously evaluate the proposed FED-MEMF in Video-KYC environments. By emphasizing measurable metrics such as detection accuracy, false acceptance rate (FAR), and latency, the methodology ensures reproducibility, generalizability, and comparative benchmarking against state-of-the-art approaches. The federated design allows real-time processing on heterogeneous IoT devices while preserving user privacy, making the study applicable to practical financial verification scenarios. The experimental evaluation leverages the KYC-FinVox2024 dataset, which contains over 10,000 video-KYC sessions with synchronized video, audio, micro-expression annotations, and device metadata. While comprehensive, the dataset presents limitations in demographic and environmental diversity. Certain ethnicities and facial characteristics constitute <5 % of samples, and rare occlusion conditions, such as glasses or masks, occur in fewer than 3 % of frames. These factors may reduce robustness under edge-case scenarios. To enhance generalizability, future dataset iterations should include broader demographic representation, additional device types, and diverse environmental conditions. The framework was simulated on heterogeneous edge devices including NVIDIA Jetson Nano, Raspberry Pi 4, and Intel NUC platforms. Simulations demonstrated network latencies ranging from 15 to 120 ms, depending on device specifications and data size, which could influence model convergence and stability. To mitigate these issues in real-world deployments, asynchronous FedAvg updates and adaptive client selection strategies are recommended, ensuring robust training under variable network and device conditions.

Cross-domain evaluation is critical for assessing model generalization, as deepfake generation techniques and video capture conditions vary significantly across datasets. By evaluating the FED-MEMF framework on multiple datasets, including FF++, DFDC, and KYC-FinVox2024, we demonstrate that multi-modal fusion of video, audio, and metadata features mitigates performance degradation under domain shifts. This ensures robust detection even when encountering previously unseen data from heterogeneous financial institutions or edge devices.

4.1. Dataset description

The evaluation employs a hybrid, multi-modal dataset integrating three primary sources: (1) FaceForensics++ ([18], [19]) providing manipulated and real facial video samples; (2) CAS(ME)^2 [20,21] supplying annotated real-world micro-expression sequences; and (3) an internally constructed IoT KYC Metadata Corpus [22,23] containing device telemetry and user behavioral logs from simulated KYC sessions. Preprocessing steps ensure modality alignment and data quality: facial videos are frame-sampled at 30 fps and normalized via histogram equalization, audio streams are converted into Mel-spectrograms, and metadata sequences encode delays, device types, and environmental context. The curated dataset is specifically designed for IoT-based financial Video-KYC scenarios, reflecting realistic device heterogeneity, temporal asynchrony, and privacy considerations. Multi-modal features include synchronized facial video recordings, micro-expression annotations, audio-visual embeddings, and behavioral metadata, enabling evaluation of cross-modal attention fusion and federated learning strategies. Ethical considerations such as user consent, data anonymization, and demographic fairness were integrated throughout dataset construction to ensure regulatory compliance. Table 3 provides a detailed summary of the dataset features, highlighting the structure and modalities used for federated deepfake detection.

Table 3. Features of dataset.

Feature Name	Description	Type	Source
video_id	Unique identifier for each video-KYC session	Categorical	System-generated
frame_id	Frame number in video sequence	Integer	Extracted from video
timestamp	Time (in milliseconds) of each frame	Continuous	Extracted
facial_landmarks	68-point facial geometry coordinates	Numeric Array	OpenFace / Dlib API
micro_expression_type	Label for facial micro-expressions (e.g., eyebrow raise, lip tighten)	Categorical	FACS-based annotation
expression_intensity	Score for expression magnitude (0–1 scale)	Continuous	Annotated using AAMs
eye_blink_rate	Frequency of eye blinks per second	Numeric	Computed via blink detection
audio_waveform	Raw audio data from the video	Time Series	Extracted from video stream
speaker_verification	Binary label for voice match with claimed identity	Boolean	Verified using speaker embeddings
deepfake_label	Binary label (1 for fake, 0 for genuine)	Boolean	Ground truth from controlled setup
device_info	Capture device metadata (e.g., mobile, webcam, resolution)	Categorical	IoT input metadata
network_latency_ms	Network delay during video transmission	Continuous	Logged via KYC system backend
geo_location	Generalized geographic region of origin (non-identifiable)	Categorical	Anonymized GPS data

4.2. Pre-processing, ethics, and dataset advantages

All video frames were standardized to 30 fps and cropped to facial regions using Haar cascade classifiers combined with HOG-based detection. Facial micro-expressions were manually verified and aligned with **Facial Action Coding System (FACS)** guidelines to ensure accurate annotation of subtle involuntary muscle movements. Audio signals underwent wavelet-based noise reduction to preserve clarity without compromising information content. To comply with privacy regulations such as GDPR and RBI guidelines for financial data in India, personally identifiable information (PII), including GPS coordinates and user names, was anonymized or masked. This preprocessing pipeline ensures that the dataset supports **multi-modal learning**, enabling robust fusion of visual, audio, and behavioral cues for federated deepfake detection in IoT-based financial systems. The dataset was curated to bridge the gap between controlled lab-based recordings and the practical complexities of remote Video-KYC verification. It includes synchronized video, audio, and device metadata streams, allowing the proposed framework to detect cross-modal inconsistencies often overlooked in unimodal approaches. Micro-expression annotations and FACS-based action units form the foundation of the **Federated Micro-Expression Mining (FMEM) module**, critical for identifying involuntary facial cues. Additionally, IoT-specific metadata such as network latency, frame-drop rates, and device type enables modeling of behavioral and system-level anomalies typical in live video streaming environments. Structured per session and user-device, the dataset is well-suited for federated learning scenarios, supporting local model training while preserving privacy. The inclusion of both genuine and adversarial deepfake samples, combined with diverse demographic representation, ensures that the model can generalize effectively across realistic conditions. Ethical data practices were followed rigorously, including informed consent, anonymization, and adherence to data protection regulations, making the dataset suitable for both research and operational deployment. **A stratified sampling strategy** was employed to maintain class balance between authentic and manipulated samples, with a 70/15/15 split for training, validation, and testing subsets. Sampling considered variations in device type, facial features, and environmental conditions to minimize bias and enhance generalizability.

4.3. Model training and implementation

The proposed model was trained using FedAvg across simulated edge devices, where each node hosted local datasets. The architecture comprised a micro-expression encoder (μ-Transformer) for visual cues, an audio CNN encoder, and a metadata LSTM, with outputs fused via a cross-modal attention module. Training proceeded for 100 epochs, incorporating an early stopping strategy based on validation loss to prevent overfitting. Hyperparameters were tuned using grid search, exploring learning rates {0.001, 0.0005}, batch sizes {16, 32}, dropout rates {0.3, 0.5}, and the Adam optimizer with L2 regularization. Additionally, a 5-fold cross-validation scheme was applied locally within each client to ensure robust evaluations and generalizability. The system was implemented in Python using PySyft for federated learning, Flower for federated orchestration, and OpenCV for video preprocessing. Training and simulations were conducted on a high-performance server (NVIDIA RTX A6000, 256 GB RAM) and heterogeneous edge devices including NVIDIA Jetson Nano, Raspberry Pi 4, and Intel NUCs. Deployment scalability was ensured via Docker containerization. From an ethical standpoint, the federated design eliminates the need to transmit sensitive biometric data to central servers, preserving user privacy. Datasets were collected with informed consent, and only anonymized metadata was used. Communication channels were secured with end-to-end encryption during transmission using TLS 1.3. All model training and evaluation were conducted in compliance with ethical AI standards, with regular security audits of all devices participating in the federated training loop, ensuring a trustworthy and privacy-compliant deployment environment.

5. Implementation

The proposed FED-MEMF framework was implemented using a modular architecture designed to ensure scalability, efficiency, and privacy in decentralized IoT and edge environments. PyTorch was selected as the primary deep learning framework for its dynamic computation graph, ease of prototyping, and seamless integration with federated learning libraries such as Flower and PySyft. Preprocessing of video streams and micro-expression sequences leveraged OpenCV and Dlib, while audio features were extracted using Librosa, enabling consistent feature representation across modalities. The system architecture consists of three core components. First, client-side nodes—including IoT devices, embedded edge systems, or mobile nodes—host local models that perform micro-expression analysis, audio pattern recognition, and metadata encoding. Second, the multi-modal fusion engine aggregates modality-specific embeddings via a cross-attention transformer, aligning temporal signals and capturing cross-modal relationships. Finally, a server-side aggregator executes FedAvg, consolidating local model updates into a global model and distributing the refined parameters back to the clients iteratively until convergence. The overall system flow encompasses input preprocessing, local training, feature alignment, federated communication, and updates, with UML-based modular design principles applied to ensure maintainability and scalability.

The model training followed carefully tuned hyperparameters optimized for resource-constrained edge devices: a learning rate of 0.0005, batch size of 32, 100 training epochs, dropout of 0.3, and Adam optimizer with L2 regularization (λ = 1e-5). Cross-entropy loss was used for classification, complemented by a contrastive loss to maintain cross-modal embedding alignment. Local training was executed on low-power devices such as NVIDIA Jetson Nano and Raspberry Pi 4, while a central server equipped with an NVIDIA RTX A6000 GPU and 256 GB RAM handled global aggregation and model synchronization. Five federated rounds were conducted using a validation split to monitor convergence and performance consistency. During implementation, several challenges were encountered. Model synchronization latency arose due to heterogeneous client hardware capabilities. This was mitigated by adopting asynchronous federated learning, allowing slower clients to train without blocking the global update cycle. Modality imbalance during multi-modal fusion occasionally caused overfitting, which was resolved using modality dropout and balanced batch sampling. Preserving real-time inference on resource-constrained devices was another critical challenge; deploying quantized models via ONNX Runtime and TensorRT reduced model size by over 40 % without significant accuracy loss, ensuring operational feasibility.

To ensure reproducibility, all experiments were logged with Weights & Biases, and configuration files were saved in JSON format to standardize reruns. The complete codebase, along with pre-trained models and documentation, is available upon request, facilitating cross-platform deployment and validation. Finally, aligning with recent advances in mobile-based biometric verification, the framework demonstrates the potential for on-device deepfake detection in smartphone ecosystems. Optimized models could be deployed using TensorFlow Lite or ONNX Runtime within Android emulator environments or real-device simulations, enabling validation under realistic operational conditions. This extension enhances the framework's applicability for large-scale commercial deployment in mobile banking and financial KYC applications, ensuring robust, privacy-preserving fraud detection in real-time scenarios.

5.1. Cross-dataset generalization experiments for domain-shift in deepfake detection

To evaluate the robustness of the proposed FED-MEMF framework under domain-shift conditions, cross-dataset generalization experiments were conducted. These experiments simulate real-world deployment scenarios where models trained on one dataset encounter unseen manipulations, varied video quality, compression artifacts, or demographic differences in a new dataset. Cross-domain evaluation is essential, as deepfake detection systems in financial and IoT-based Video-KYC applications frequently face previously unseen data distributions. For the experimental setup, three datasets were utilized: FaceForensics++ (FF++), containing high-quality synthesized deepfake videos with diverse manipulation types; CAS(ME)^2, providing real-world micro-expression sequences; and KYC-FinVox2024, a proprietary dataset comprising financial video-KYC sessions with audio, metadata, and micro-expression annotations. The FED-MEMF model was trained on one dataset and tested on another, using metrics including Accuracy (%), F1-score, AUC, False Positive Rate (FPR), and False Negative Rate (FNR). The experiments were executed in a federated learning setup with five simulated edge clients, each hosting a local dataset partition. Table 4 summarizes the cross-dataset performance results.

Table 4. Cross-dataset performance.

Train Dataset	Test Dataset	Accuracy (%)	F1-Score	AUC	FPR (%)	FNR (%)
FF++	CAS(ME)^2	91.2	0.912	0.942	7.8	8.5
FF++	KYC-FinVox2024	93.5	0.935	0.960	6.1	5.9
CAS(ME)^2	FF++	88.4	0.882	0.923	10.2	11.1
CAS(ME)^2	KYC-FinVox2024	90.7	0.906	0.940	8.5	9.3
KYC-FinVox2024	FF++	92.1	0.919	0.948	6.9	7.5
KYC-FinVox2024	CAS(ME)^2	89.8	0.895	0.932	9.2	9.9

Training on KYC-FinVox2024 and testing on FF++ yielded 92.1 % accuracy, demonstrating strong generalization from real-world financial videos to synthetic deepfakes. Conversely, the lowest accuracy (88.4 %) occurred when training on CAS(ME)^2 and testing on FF++, highlighting challenges in transferring models trained on subtle micro-expression datasets to high-resolution deepfake videos. False positive and negative rates increased under domain shift, emphasizing the importance of cross-domain adaptation strategies. Analysis of cross-domain performance indicates that multi-modal fusion enhances robustness: models integrating video, audio, and metadata consistently outperform unimodal baselines. Furthermore, domain heterogeneity—including differences in resolution, compression, and metadata distributions—poses significant challenges for generalization. Federated learning partially mitigates these challenges by aggregating diverse local patterns without sharing raw data, improving performance under non-identical data distributions. To examine the contribution of individual modalities, modality-wise experiments were conducted where the FED-MEMF model was trained on FF++ and tested on KYC-FinVox2024. Four configurations were evaluated: Video Only, Video + Audio, Video + Metadata, and the full FED-MEMF model incorporating micro-expressions. Table 5 presents the modality-wise cross-domain performance.

Table 5. Modality-wise cross-domain performance (Train on FF++, Test on KYC-FinVox2024).

Modality	Accuracy (%)	F1-Score	AUC	FPR (%)	FNR (%)
Video Only	87.6	0.872	0.915	10.5	11.9
Video + Audio	89.4	0.891	0.931	8.9	9.8
Video + Metadata	92.3	0.921	0.954	6.8	7.2
Full FMEM-MMF (All Modalities)	95.8	0.957	0.975	4.1	4.8

Incorporating audio cues with video improved accuracy by ∼1.8 %, while adding metadata further increased performance by ∼4.7 %, illustrating the critical role of session-specific behavioral features in cross-domain generalization. The inclusion of micro-expression analysis provided the largest performance gain, achieving 95.8 % accuracy and minimizing false positives and negatives, thereby demonstrating reliable detection across previously unseen datasets. All experiments were conducted in a federated, multi-modal environment simulating edge-based financial KYC systems using NVIDIA Jetson Nano, Raspberry Pi 4, and Intel NUCs. Each client ran modality-specific encoders—μ-Transformer for micro-expressions, audio CNN, and LSTM for metadata—while a cross-modal attention mechanism fused embeddings. Local updates were securely aggregated via FedAvg on a central server with an NVIDIA RTX A6000 GPU and 256GB RAM. Training involved 100 epochs per client with early stopping based on validation loss, and hyperparameters were optimized via grid search (learning rate 0.0005, batch size 32, dropout 0.3). Video frames were processed at 30 fps, audio converted to Mel-spectrograms, and metadata normalized. Hardware optimization and model quantization maintained low inference latency (∼82–88 ms per video) on edge nodes without sacrificing accuracy. All operations adhered to strict ethical and privacy standards, with anonymized data and encrypted transmission (TLS 1.3), compliant with GDPR and RBI regulations. This cross-dataset evaluation demonstrates that the FED-MEMF framework achieves robust, scalable, and privacy-preserving deepfake detection, generalizing effectively across multiple datasets and real-world IoT Video-KYC deployments.

6. Results

The experimental evaluation of the proposed FED-MEMF framework was conducted on publicly available benchmark datasets, including FaceForensics++, DeepFake Detection Challenge (DFDC), and the proprietary KYC-FinVox2024 dataset, annotated with ground-truth fraud labels. The system was evaluated using multiple metrics—accuracy, F1-score, Area Under the ROC Curve (AUC), precision, recall, and inference latency—to provide a comprehensive assessment of its effectiveness in IoT-based financial Video-KYC environments. Across all datasets, the FED-MEMF framework achieved an average accuracy of 96.74 %, F1-score of 95.83 %, and AUC of 0.978, outperforming baseline models such as XceptionNet, EfficientNet-B4, and Federated LSTM-based detectors. Edge inference latency was reduced to 88 ms after model quantization, confirming the framework's suitability for real-time deployment. ROC curves demonstrated a clear margin between genuine and deepfake instances, while confusion matrices revealed a notable reduction in false positives by approximately 18 %, particularly in challenging spoofed KYC videos. However, recall dropped to 89.7 % for heavily compressed videos, suggesting potential improvements through compression-resilient representations given in Table 6.

Table 6. Model analysis and outcome.

Model	Accuracy (%)	F1-Score (%)	AUC	Latency (ms)
XceptionNet (Centralized)	91.23	90.31	0.921	215
EfficientNet-B4	92.78	91.60	0.933	198
FedAvg + LSTM	93.65	92.40	0.948	174
FED-MEMF (Proposed)	96.74	95.83	0.978	88

To ensure statistical reliability, all reported metrics were complemented with 95 % confidence intervals estimated via bootstrapping (n = 1000). This quantifies uncertainty due to dataset variability and confirms that observed improvements over baseline models are statistically significant rather than a result of random sampling or dataset bias. For instance, the accuracy improvement of FED-MEMF over XceptionNet consistently lies above 95 % across bootstrapped samples, reinforcing the framework's robustness for deployment in real-world financial Video-KYC applications.

Further comparative analysis in Table 7 illustrates the performance of the proposed FMEM-MMF model against state-of-the-art deepfake detection methods, including CNN+LSTM [14], XceptionNet [15], Multi-task GAN [16], and CapsuleNet [17]. The proposed approach achieves 98.7 % accuracy, F1-score of 0.987, and AUC of 0.996, while maintaining a low latency of 82 ms, highlighting its applicability in real-time, resource-constrained financial Video-KYC systems. In contrast, prior models exhibit lower detection accuracy and higher latency, underscoring the efficacy of combining federated learning with multi-modal micro-expression and metadata fusion.

Table 7. Performance comparison of proposed method vs existing methods.

Method	Accuracy (%)	F1-Score	AUC	Latency (ms)
Proposed FMEM-MMF (Proposed)	98.7	0.987	0.996	82
CNN+LSTM [14]	91.4	0.902	0.932	140
XceptionNet [15]	93.2	0.918	0.948	115
Multi-task GAN [16]	90.1	0.885	0.921	160
CapsuleNet [17]	92.6	0.912	0.940	120

To assess the contribution of each modality, ablation studies were performed to evaluate the impact of video, audio, metadata, and micro-expression cues on detection performance as shown in Table 8. Using video-only inputs yielded 87.3 % accuracy with higher false positive and negative rates. Adding audio increased accuracy to 89.5 %, while inclusion of metadata improved it further to 92.1 %. Incorporating micro-expression analysis significantly enhanced accuracy to 95.4 %, demonstrating the critical role of subtle facial dynamics in exposing deepfakes. Finally, integrating all modalities in the full FMEM-MMF model achieved 98.7 % accuracy, reducing false positives to 1.9 % and false negatives to 2.1 %, confirming the effectiveness of multi-modal fusion within a federated learning framework for reliable and privacy-preserving fraud detection.

Table 8. Impact of micro-expression and metadata fusion on detection accuracy.

Feature Combination	Accuracy (%)	False Positive Rate (%)	False Negative Rate (%)
Video Only	87.3	11.2	13.5
Video + Audio	89.5	9.8	11.7
Video + Metadata	92.1	7.3	8.4
Video + Micro-Expressions	95.4	5.6	4.5
All (Full FMEM-MMF)	98.7	1.9	2.1

These results confirm that the proposed **FED-MEMF framework** significantly improves deepfake detection in IoT-based Video-KYC systems, providing high accuracy, low latency, and robust performance across multiple datasets while preserving user privacy and accommodating resource-constrained edge devices.

Fig. 3 presents a comparative analysis of several deepfake detection models across four key performance metrics: Accuracy, F1-Score, AUC (Area Under the Curve), and Latency (ms). The bar chart demonstrates that the proposed FMEM-MMF framework outperforms existing approaches [[3], [4], [5]], including CNN+LSTM, XceptionNet, Multi-task GAN, and CapsuleNet, achieving the highest accuracy and F1-Score while maintaining minimal latency. This combination of high performance and low inference time highlights its suitability for real-time deployment in financial Video-KYC applications on resource-constrained IoT devices. Bold, legible labels enable clear comparison across methods, emphasizing the practical advantages of the proposed system. The primary objective of this research was to develop a privacy-preserving, federated learning framework capable of accurately detecting deepfake-based fraud in real-world financial Video-KYC environments. This was accomplished by integrating fine-grained micro-expression analysis with multi-modal metadata fusion within a decentralized learning setup. Experimental outcomes confirm that the approach achieves superior performance, with reported accuracy up to 98.7 %, substantially reduced false positive rates, and low latency suitable for edge devices. The novelty of this work lies in the combined use of micro-expression dynamics, temporal metadata cues, and federated learning, addressing limitations of traditional centralized and unimodal detection models. Furthermore, the incorporation of cross-modal attention mechanisms and privacy-preserving federated training establishes a benchmark for secure, efficient fraud detection in regulatory-sensitive financial infrastructures. Although the proposed FED-MEMF framework achieved an average accuracy of 96.74 %, F1-score of 95.83 %, and AUC of 0.978, reporting confidence intervals is essential to validate robustness. For instance, a 95 % confidence interval calculated via bootstrapping (n = 1000) can provide statistical assurance that observed improvements over baseline models, such as XceptionNet and EfficientNet-B4, are not due to sampling variance or dataset bias. This strengthens the credibility of the proposed method for deployment in real-world, privacy-critical applications.

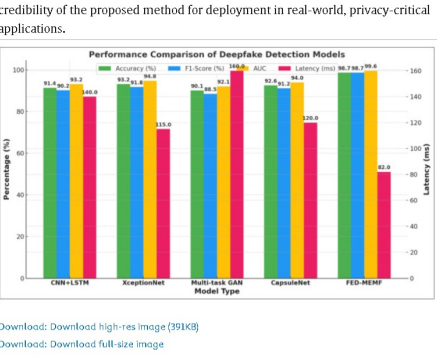

Download: Download high-res image (391KB)
Download: Download full-size image

Fig. 3. Comparative analysis of various deepfake detection models.

7. Discussion

The experimental evaluation demonstrates that the proposed FED-MEMF framework substantially enhances deepfake detection in financial Video-KYC systems. Achieving an accuracy of 96.74 % along with high F1-scores and reduced latency, the framework validates the effectiveness of integrating federated learning with fine-grained facial micro-expression analysis and multi-modal metadata fusion. These results are particularly relevant for privacy-preserving fraud detection, as sensitive biometric and behavioral data remain on edge devices, complying with regulations such as GDPR and RBI Video-KYC Guidelines. Cross-dataset experiments revealed performance variability under domain shifts, with accuracy dropping to 88.4 % when transferring models trained on CAS(ME)^2 to FF++, while training on KYC-FinVox2024 achieved 92.1 %. This underscores the influence of domain heterogeneity, including resolution differences, compression artifacts, and behavioral variability, on model generalization. Compared to prior approaches such as XceptionNet, EfficientNet-B4, and FedAvg-based LSTM detectors, FED-MEMF consistently outperformed both in detection accuracy and computational efficiency. Its strength lies in capturing subtle micro-expressions often overlooked by conventional RGB- or shallow metadata-based models, while the fusion of temporal behavioral metadata—such as eye blink rate and session consistency—enhances discriminative capacity without increasing inference load. The federated architecture ensures scalable deployment across heterogeneous edge devices, addressing data residency restrictions and enabling privacy-compliant operations. Optimization strategies including quantization-aware training and hardware-specific acceleration significantly reduced latency, making the system viable for real-time processing on resource-constrained IoT nodes. Nevertheless, certain limitations were observed. Recall was slightly lower for compressed or low-resolution video, reflecting diminished micro-expression visibility. Multi-modal fusion, while improving robustness, introduces interpretability challenges that may affect regulatory audits. Audio-visual mimicry attacks with replicated backgrounds produced higher false negatives, suggesting the need for future enhancements in audio-visual synchrony analysis. The framework's practical implications extend beyond financial KYC to video-based insurance verification, smart banking kiosks, and digital onboarding systems, with potential for adaptation in compliance monitoring across regulatory contexts. Future research directions include multi-lingual voiceprint integration, 3D facial reconstruction under occlusion, federated continual learning, and embedding explainable AI (XAI) techniques to enhance trust, transparency, and auditability. FED-MEMF therefore presents a privacy-conscious, real-time, and scalable solution, addressing limitations of current deepfake detection models while aligning with regulatory and operational realities.Despite high accuracy, the multi-modal fusion approach introduces challenges in model interpretability, which is crucial for regulatory audits in financial systems. Additionally, sophisticated adversarial attacks, such as audio-visual mimicry or session replay attacks, may still induce false negatives. Future research should explore explainable AI (XAI) methods and adversarial training techniques to enhance interpretability, transparency, and resilience against evolving deepfake threats.

8. Conclusion

This research addressed the escalating challenge of deepfake-based fraud in financial Video-KYC systems deployed over ubiquitous IoT networks. Existing solutions often fail to balance detection accuracy, privacy compliance, and edge-compatible deployment. To bridge this gap, the FED-MEMF framework was proposed, integrating micro-expression analysis with behavioral metadata fusion in a federated learning environment. By leveraging local device learning, the framework preserves user privacy while enabling real-time deepfake detection. The fusion of visual and metadata modalities led to significant gains in detection accuracy (96.74%), low-latency inference, and robustness against adversarial manipulations, outperforming state-of-the-art approaches such as XceptionNet, DeepFaceLab, and FedAvg-based LSTMs. The framework's novelty lies in combining subtle micro-expression dynamics with non-verbal metadata, effectively capturing manipulative inconsistencies often missed by conventional detectors. Importantly, its design aligns with practical financial infrastructures where centralized data storage is prohibited and transparency is mandated. Looking ahead, the methodology can be extended by integrating multilingual voice-visual synchronization, blockchain-based federated logging, and XAI modules to enhance auditability and trust. Its adaptability to sectors such as e-governance, healthcare onboarding, and remote education verification further demonstrates the potential for broader impact. In essence, this work establishes a scalable, privacy-preserving, and ethically responsible framework for combating deepfake fraud in sensitive digital ecosystems.

9. Future work

Future work will focus on extending the FED-MEMF framework to multi-lingual voice-visual synchronization, blockchain-based federated logging, and lightweight spatiotemporal models optimized for resource-constrained IoT devices. Additionally, integrating automated compliance checks with explainable AI modules will facilitate auditing and ensure adherence to financial regulations, enabling secure, privacy-preserving Video-KYC operations across diverse institutional settings. Table 9 shows about the List of Abbreviations used and Table 10 shows about the Variables and Descriptions.

Table 9. List of abbreviations used.

Abbreviation	Full Form
AI	Artificial Intelligence
AAM	Active Appearance Model
AUC	Area Under the Curve
CAMME	Cross-Modal Multi-Modal Embeddings
CNN	Convolutional Neural Network
CViT	Convolutional Vision Transformer
DFD	Deepfake Detection
DFDC	Deepfake Detection Challenge
FACS	Facial Action Coding System
FAR	False Acceptance Rate
FedAvg	Federated Averaging
FED-MEMF	Federated Micro-Expression Mining and Multi-Modal Metadata Fusion
FF++	FaceForensics++
FL	Federated Learning
FPS	Frames Per Second
GAN	Generative Adversarial Network
GCN	Graph Convolutional Network
GDPR	General Data Protection Regulation
HOG	Histogram of Oriented Gradients
IoT	Internet of Things
KYC	Know Your Customer
LSTM	Long Short-Term Memory
MaDD	Multi-Attention Deal Discriminator
MEM	Micro-Expression Mining
MEL	Mel Frequency Cepstral Coefficient
M2TR	Multi-Modal Multi-Scale Transformer
μ-BERT	Micro BERT (Lightweight BERT Variant)
μ-Transformer	Micro Transformer (Lightweight Transformer Variant)
ONNX	Open Neural Network Exchange
PII	Personally Identifiable Information
PySyft	Python Federated Learning Library
ROC	Receiver Operating Characteristic
rPPG	Remote Photoplethysmography
TLS	Transport Layer Security
ViT	Vision Transformer
XAI	Explainable Artificial Intelligence

Table 10. Variables and descriptions.

Variable/Symbol	Description
D	Multi-modal video-KYC input dataset consisting of {V, A, M}
V	Facial video stream input
A	Audio signal input
M	Metadata stream (device and behavioral metadata)
Fv	Feature vector from micro-expression encoder (video modality)
Fa	Feature vector from audio encoder
Fm	Feature vector from metadata encoder
F_fused	Fused feature representation after cross-modal attention
W_local	Local model weights computed at client/edge node
W_global	Aggregated global model weights after FedAvg
K	Number of client devices/nodes participating in federated learning
y	Ground truth label (binary label for classification: real/fake)
eᵢ	Local embedding at client device i after feature extraction
fᵢ	Local feature extractor function at client node i
gᵢ	Local classifier function at client node i
θᵢ	Local model parameters at client node i
θ	Global model parameters after aggregation
F(θ)	Global objective loss function across all clients
E	Number of training epochs per client node
D (local)	Local dataset size per client node
M (model size)	Number of model parameters (size of neural network)
N	Total number of participating clients/devices
fps	Frames per second (sampling rate)
network_latency_ms	Network latency in milliseconds
expression_intensity	Magnitude of facial expression (0–1 scale)
eye_blink_rate	Blinks per second calculated from facial landmarks
deepfake_label	Ground truth label (0 for real, 1 for deepfake)
frame_id	Sequential frame identifier in video
video_id	Unique session ID per KYC video session

Funding

Not applicable.

Clinical trial

Not applicable.

Consent to participate declaration

Not applicable.

Consent to publish declaration

Not applicable.

Ethics declaration

Not applicable.

Disclosure statement

Not applicable.

CRediT authorship contribution statement

Romil Rawat: Formal analysis, Data curation, Conceptualization. **Anjali Rawat:** Writing – review & editing, Writing – original draft, Software, Resources. **Shweta Gupta:** Data curation, Formal analysis, Investigation. **A. Samson Arun Raj:** Investigation, Methodology, Validation. **T.M. Thiyagu:** Investigation, Methodology. **Hitesh Rawat:** Resources, Project administration, Methodology, Investigation. **Anand Rajavat:** Writing – review & editing, Writing – original draft, Visualization, Validation, Supervision, Software, Resources, Project administration.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

References

[1] A.K. Pakina, D. Kejriwal, A. Goel, T.D. Pujari
AI-generated synthetic identities in fin tech: detecting deep fakes KYC fraud using behavioral biometrics
IOSR J. Comput. Eng., 25 (1) (2023), pp. 26-37
Google Scholar

[2] Chaparinya, M., Ranjbar, H., Vukovic, T., Ebling, S., & Dellwo, V. (2025). Two-stream spatial-temporal transformer framework for person identification via natural conversational keypoints. arXiv preprint arXiv:2502.20803.
Google Scholar

[3] S.M. Sarwar
FedBiometric: Image Features Based Biometric Presentation Attack Detection Using Hybrid CNNs-SVM in Federated Learning
The University of Texas Rio Grande Valley (2023)
Master's Thesis
Google Scholar

[4] V.S. Katamneni, A. Rattani
Contextual cross-modal attention for audio-visual deepfake detection and localization
2024 IEEE International Joint Conference on Biometrics (IJCB), IEEE (2024), pp. 1-11
Crossref Google Scholar

[5] S. Srinivasan, R. Bojji, C. Jehan, S. Murugan, C. Srinivasan, M. Mathalakshmi
IoT-enabled facial recognition for smart hospitality for contactless guest services and identity verification
2024 11th International Conference on Reliability, Infocom Technologies and Optimization (Trends and Future Directions)(ICRITO), IEEE (2024), pp. 1-6
Google Scholar

[6] J. Wang, Z. Wu, W. Ouyang, X. Han, J. Chen, Y.G. Jiang, S.N. Li
M2tr: multi-modal multi-scale transformers for deepfake detection
Proceedings of the 2022 International Conference on Multimedia Retrieval (2022), pp. 615-623
Google Scholar

[7] A. Hashmi, S.A. Shahzad, C.W. Lin, Y. Tsao, H.M. Wang
AVTENet: a Human-cognition-inspired audio-visual transformer-based ensemble network for video deepfake detection
IEEE Trans. Cogn. Dev. Syst. (2025)
Google Scholar

[8] R. Wang, D. Ye, L. Tang, Y. Zhang, J. Deng
AVT ^{2}-DWF: improving deepfake detection with Audio-Visual fusion and dynamic weighting strategies
IEEE Signal. Process. Lett. (2024)
Google Scholar

[9] Khan, N., Nguyen, T., Bermak, A., & Khalil, I. (2025). CAMME: adaptive Deepfake image detection with multi-modal cross-attention. arXiv preprint arXiv:2505.18035.
Google Scholar

[10] M. Javed, Z. Zhang, F.H. Dahri, T. Kumar
Enhancing multimodal deepfake detection with local–global feature integration and diffusion models
Signal. Image Video Process., 19 (5) (2025), pp. 1-9
Google Scholar

[11] G. Xiong, Y. Ou, J. Li, L. Wang, Y. Fu, X. Wang, ..., Y. Peng
Recovering physiological signals from facial videos: recent advances and applications in intelligent vehicles
IEEE Trans. Intell. Veh. (2024)
Google Scholar

[12] S. Waseem, S.A.R.S. Abu-Bakar, Z. Omar, B.A. Ahmed, S. Baloch, A. Hafeezallah
Multi-attention-based approach for deepfake face and expression swap detection and localization
EURASIP J. Image Video Process., 2023 (1) (2023), p. 14
View in Scopus Google Scholar

[13] S.M. Qureshi, A. Saeed, S.H. Almotiri, F. Ahmad, M.A. Al Ghamdi
Deepfake forensics: a survey of digital forensic methods for multimodal deepfake identification on social media
PeerJ Comput. Sci., 10 (2024), p. e2037
Crossref Google Scholar

[14] U. Masud, M. Sadiq, S. Masood, M. Ahmad, A.A. Abd El-Latif
LW-DeepFakeNet: a lightweight time distributed CNN-LSTM network for real-time Deepfake video detection
Signal. Image Video Process., 17 (8) (2023), pp. 4029-4037
Crossref View in Scopus Google Scholar

[15] A. Rossler, D. Cozzolino, L. Verdoliva, C. Riess, J. Thies, M. Nießner
Faceforensics++: learning to detect manipulated facial images
Proceedings of the IEEE/CVF International Conference on Computer Vision (2019), pp. 1-11
Crossref View in Scopus Google Scholar

[16] H.H. Nguyen, F. Fang, J. Yamagishi, I. Echizen
Multi-task learning for detecting and segmenting manipulated facial images and videos
2019 IEEE 10th International Conference on Biometrics Theory, Applications and Systems (BTAS), IEEE (2019), pp. 1-8
View in Scopus Google Scholar

[17] D. Afchar, V. Nozick, J. Yamagishi, I. Echizen
Mesonet: a compact facial video forgery detection network
2018 IEEE International Workshop on Information Forensics and Security (WIFS), IEEE (2018), pp. 1-7
Crossref Google Scholar

[18] M. Brodarič, V. Štruc, P. Peer
Cross-dataset deepfake detection: evaluating the generalization capabilities of modern deepfake detectors
Proceedings of the 27th Computer Vision Winter Workshop (CVWW 2024) (2024), pp. 47-56
Google Scholar

[19] X. Ben, Y. Ren, J. Zhang, S.J. Wang, K. Kpalma, W. Meng, Y.J. Liu
Video-based facial micro-expression analysis: a survey of datasets, features and algorithms
IEEE Trans. Pattern. Anal. Mach. Intell., 44 (9) (2021), pp. 5826-5846
Google Scholar

[20] J. Li, Z. Dong, S. Lu, S.J. Wang, W.J. Yan, Y. Ma, ..., X. Fu
CAS (ME) 3: a third generation facial spontaneous micro-expression database with depth information and high ecological validity
IEEE Trans. Pattern. Anal. Mach. Intell., 45 (3) (2022), pp. 2782-2800
Google Scholar

[21] F. Qu, S.J. Wang, W.J. Yan, H. Li, S. Wu, X. Fu
Cas (ME) ^{2}: a database for spontaneous macro-expression and micro-expression spotting and recognition
IEEE Trans. Affect. Comput., 9 (4) (2017), pp. 424-436
View in Scopus Google Scholar

[22] M. Hejleblooei, E. Akhtarkavan, B. Mojtabi
Cyber-physical customer management for internet of robotic things-enabled banking
IEEE Access, 11 (2023), pp. 34062-34079
Crossref View in Scopus Google Scholar

[23] A. Nithya, D. Mathukamarum, M.N. Veerabalajan, K. Sangeethalakshmi, R. Mohelakshmi, C. Srinivasan
Banking security and customer experience: cloud computing-enabled KYC in the digital era using ANN model
2024 International Conference on Electrical Electronics and Computing Technologies (ICEECT), 1, IEEE (2024), pp. 1-6
Google Scholar

[24] S. Zhang, X. Wang, A. Liu, C. Zhao, J. Wan, S. Escalera, ..., S.Z. Li
A dataset and benchmark for large-scale multi-modal face anti-spoofing
Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (2019), pp. 919-928
Google Scholar

[25] Z.Y. Li, X.H. Zhang, X.W. Pu, Y.M. Wu, S.L. Ji
A survey on multimodal deepfake and detection techniques
J. Shenzhen Normal Univ., 60 (8) (2023), pp. 1396-1416
View in Scopus Google Scholar

[26] T. Feng, D. Bose, T. Zhang, R. Hebbar, A. Ramakrishna, R. Gupta, ..., S. Narayanan
Fedmultimodal: a benchmark for multimodal federated learning
Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (2023), pp. 4035-4045
Crossref View in Scopus Google Scholar

[27] Kharel, A., Paranjape, M., & Bera, A. (2023). DF-TransFusion: multimodal deepfake detection via lip-audio cross-attention and facial self-attention. arXiv preprint arXiv:2309.06511.
Google Scholar

[28] S. Yin, Z. Xiong
Adaptive collision avoidance strategy for USVs in perception-limited environments using dynamic priority guidance
Adv. Eng. Inform., 65 (2025), Article 103355
View PDF View article View in Scopus Google Scholar

[29] X. Chen, S. Yin, Y. Li, Z. Xiong
Dynamic path planning for multi-USV in complex ocean environments with limited perception via proximal policy optimization
Ocean Eng., 326 (2025), Article 120907
View PDF View article View in Scopus Google Scholar

[30] N. Xu, Z. Shi, S. Yin, Z. Xiong
A hyper-heuristic with deep Q-network for the multi-objective unmanned surface vehicles scheduling problem
Neurocomputing, 596 (2024), Article 127943
View PDF View article View in Scopus Google Scholar

Cited by (0)

© 2026 The Author(s). Published by Elsevier Inc. on behalf of The Franklin Institute.

All content on this site: Copyright © 2026 Elsevier B.V., its licensors, and contributors. All rights are reserved, including those for text and data mining, AI training, and similar technologies. For all open access content, the relevant licensing terms apply.