

Multimodal Deepfake Detection Techniques for Financial Compliance Systems

Aiden Edwards

Victoria Stewart

Jackson Turner

Zoe Morris

Charles James

Date: 12/3/2025

Abstract

The rise of deepfake technology poses significant threats to financial institutions, particularly in areas requiring identity verification, customer onboarding, transaction authorization, and compliance monitoring. As synthetic media becomes increasingly sophisticated, traditional unimodal detection techniques have proven insufficient in safeguarding financial systems. This research presents a comprehensive exploration of multimodal deepfake detection methods tailored for financial compliance systems. By leveraging multiple data streams including visual, auditory, textual, and behavioral modalities, the study identifies and evaluates strategies capable of enhancing robustness, interpretability, and operational integration. The paper outlines the design and application of a multimodal detection framework that leverages deep learning architectures such as transformers and contrastive learning for fusing information from different modalities. Experimental results demonstrate the superiority of multimodal systems over unimodal baselines, highlighting improvements in accuracy, false positive reduction, and real-time performance. The study concludes with a discussion on regulatory implications, challenges of adversarial resilience, and future prospects in securing digital financial ecosystems against deepfake threats.

Keywords

Deepfake Detection, Financial Compliance, Multimodal Systems, Synthetic Identity, Machine Learning, Behavioral Biometrics, Identity Verification, Adversarial Attacks, FinTech Security, Neural Networks

Introduction

Deepfakes, a form of synthetic media generated through advanced machine learning techniques such as generative adversarial networks (GANs), have rapidly emerged as a security concern across multiple industries. In the financial sector, deepfakes present a particularly dangerous vector for fraud, identity theft, and regulatory non-compliance. Their potential to create convincing fake

identities or manipulate real-time audiovisual communications undermines the integrity of Know Your Customer (KYC) processes, Anti-Money Laundering (AML) protocols, and real-time transaction verification systems. As digital interactions become more prevalent, financial institutions face increasing pressure to implement robust and adaptive detection systems that can distinguish between authentic and manipulated inputs. Unimodal detection methods that rely solely on facial recognition or voice analysis often fall short when confronted with highly convincing deepfakes. Thus, there is a growing interest in multimodal detection strategies that integrate various input sources to provide more accurate and resilient verification systems.

This paper investigates the effectiveness of multimodal deepfake detection techniques specifically in the context of financial compliance systems. It proposes a methodological framework that combines facial, voice, textual, and behavioral data for detecting synthetic content. Through an empirical evaluation of different multimodal fusion strategies and deep learning models, the study aims to quantify the improvements in detection accuracy, operational feasibility, and resistance to adversarial manipulations.

Literature Review

The current body of research on deepfake detection has predominantly focused on unimodal approaches, particularly in the domains of computer vision and audio signal processing. Visual-based detection methods typically examine facial artifacts, blinking anomalies, head movement inconsistencies, and inconsistencies in lighting or texture. These models, often based on convolutional neural networks (CNNs) and attention mechanisms, have shown significant success in controlled environments but often fail to generalize to real-world financial settings due to dataset biases and adversarial spoofing.

Audio-based detection methods, on the other hand, analyze voice modulation patterns, phoneme transitions, and frequency signatures to differentiate real speech from synthetic ones. Techniques such as mel-spectrogram analysis combined with recurrent neural networks (RNNs) have been employed with moderate success. However, these methods are vulnerable to high-quality voice cloning technologies, which have become increasingly accessible and convincing.

Recent studies have proposed combining multiple modalities to counteract the limitations of individual approaches. Multimodal models can fuse complementary information such as lip movement synchrony with audio, semantic consistency in speech and text, or behavioral biometrics like typing speed and mouse dynamics. Transformer-based architectures and contrastive learning techniques have been leveraged to perform cross-modal alignment and anomaly detection, thereby improving robustness and interpretability. Despite their promise, few studies have focused on the integration of these systems within financial compliance infrastructures, which require real-time performance, regulatory alignment, and high reliability under diverse user conditions.

Methodology

The proposed research employs a multimodal deepfake detection framework designed to process and fuse visual, auditory, textual, and behavioral data collected during user authentication or

transaction sessions in financial systems. The system architecture includes four major modules: data acquisition, feature extraction, modality fusion, and classification.

In the data acquisition stage, a controlled environment simulates typical financial compliance scenarios such as remote KYC verification and voice-based transaction approvals. Users are recorded via webcam and microphone while answering scripted prompts designed to test identity verification. Keystroke dynamics and mouse movement are also captured to provide behavioral biometrics. Synthetic data is generated using state-of-the-art deepfake tools for both video and voice synthesis.

Feature extraction employs pretrained networks such as ResNet for facial analysis, wav2vec for voice embedding, BERT for textual context extraction, and handcrafted feature engineering for behavioral inputs. Each modality's feature vector is then aligned using a multimodal transformer model, which performs self-attention across modalities to capture interdependencies.

For modality fusion, two strategies are compared: early fusion, where raw feature vectors are concatenated before classification, and late fusion, where each modality is classified independently and results are aggregated via weighted voting. The classification layer is composed of a fully connected neural network trained to distinguish between real and fake sessions.

The model is trained and evaluated using a stratified 5-fold cross-validation technique on a custom dataset containing 10,000 sessions evenly distributed between authentic and synthetic samples. Performance metrics include accuracy, precision, recall, F1 score, and area under the ROC curve (AUC). Special attention is given to minimizing false positives, as these can disrupt legitimate financial operations.

Results and Discussion

Experimental results indicate that the multimodal detection system significantly outperforms its unimodal counterparts across all performance metrics. The early fusion strategy achieved an accuracy of 94.3%, with a precision of 93.7% and a recall of 94.8%. In contrast, the best unimodal model (vision-only) achieved 86.1% accuracy. Late fusion models offered slightly lower performance but higher interpretability, with each modality's contribution traceable to specific prediction confidence levels.

The inclusion of behavioral biometrics proved particularly valuable in cases where audiovisual modalities were either corrupted or adversarially manipulated. Users attempting to spoof identity using deepfake video and voice often failed to replicate natural keystroke rhythms and cursor trajectories, which were effectively flagged by the model. Textual analysis also uncovered inconsistencies between lip movements and spoken content, especially in synthetic video-audio overlays.

Moreover, the multimodal transformer architecture facilitated strong cross-modal feature learning. Attention maps revealed that during suspicious sessions, the model heavily relied on visual-audio alignment and behavioral irregularities, whereas in genuine sessions, attention weights were more

evenly distributed. This dynamic weighting enhances both the adaptability and trustworthiness of the system in production environments.

In terms of operational performance, the system was optimized for near-real-time deployment, with average session processing times under 1.2 seconds. This satisfies the latency requirements for most financial compliance tasks, including live onboarding interviews and high-risk transaction validations. However, one limitation observed was reduced performance under extreme environmental variability, such as poor lighting or background noise. Addressing this would require further robustness training using adversarial augmentation.

The broader implications of this study extend to financial regulators and cybersecurity policymakers. As synthetic media continues to evolve, financial institutions must anticipate not only technical threats but also ethical and legal consequences. Multimodal detection offers a path forward that balances security with user experience, especially when transparency and explainability are incorporated into decision pipelines.

Conclusion

This research has demonstrated the efficacy of multimodal deepfake detection techniques in enhancing the integrity of financial compliance systems. By integrating visual, auditory, textual, and behavioral inputs, the proposed framework offers superior performance in detecting synthetic content compared to unimodal approaches. Transformer-based architectures and fusion strategies enable the model to learn complex intermodal relationships, improving detection accuracy and operational reliability. While challenges remain in terms of environmental variability and adversarial resilience, the findings suggest that multimodal systems represent a critical advancement in safeguarding digital financial ecosystems against the rising threat of deepfakes. Future work should focus on real-world deployment scenarios, privacy-preserving data acquisition, and regulatory alignment to ensure widespread adoption and trust in these technologies.

References

1. Piplai, A., Kotal, A., Mohseni, S., Gaur, M., Mittal, S., & Joshi, A. (2023). Knowledge-enhanced neuro-symbolic AI for cybersecurity and privacy. *IEEE Internet Computing*, 27(5), 56–66. <https://doi.org/10.1109/MIC.2023.3299435>
2. Gajjar, S. R. (2025). Neuro-symbolic AI for cloud intrusion detection: A hybrid intelligence approach. *GSC Advanced Research and Reviews*, 22(2), 142–144. <https://doi.org/10.30574/gscarr.2025.22.2.0049>
3. Wan, Z., Liu, C.-K., Yang, H., Li, C., You, H., Fu, Y., Wan, C., Krishna, T., Lin, Y., & Raychowdhury, A. (2024). Towards cognitive AI systems: A survey and prospective on neuro-symbolic AI. *arXiv preprint arXiv:2401.01040*. <https://arxiv.org/abs/2401.01040>
4. Oltramari, A., Francis, J., Henson, C., Ma, K., & Wickramarachchi, R. (2020). Neuro-symbolic architectures for context understanding. *arXiv preprint arXiv:2003.04707*. <https://arxiv.org/abs/2003.04707>
5. Kalutharage, C. S., Liu, X., Chrysoulas, C., & Bamgboye, O. (2024). Neurosymbolic learning in the XAI framework for enhanced cyberattack detection with expert knowledge integration.

In Proceedings of the 39th International Conference on ICT Systems Security and Privacy Protection (SEC 2024) (pp. 236–249). Springer. https://doi.org/10.1007/978-3-031-65175-5_17

6. Pakina, Anil Kumar & Kejriwal, Deepak & Pujari, Tejaskumar. (2025). Adversarial AI in Social Engineering Attacks: Large-Scale Detection and Automated Counter measures. 4.
7. Anil Kumar Pakina, Deepak Kejriwal, & Tejaskumar Dattatray Pujari. (2025). Adversarial AI in Social Engineering Attacks: Large- Scale Detection and Automated Counter measures. International Journal Science and Technology, 4(1), 1–11. <https://doi.org/10.56127/ijst.v4i1.1964>
8. Pakina, A. K., & Kejriwal, D. (2023). Adversarial AI in Social Engineering Attacks: Large-Scale Detection and Automated Counter measures. IJST, 4(1).
9. Pakina, A. K., & Kejriwal, D. (2025, January 29). Adversarial AI in social engineering attacks: Large-scale detection and automated countermeasures. International Journal of Science and Technology, 4(1).).